

ediction-using-arma-end-to-end

April 16, 2024

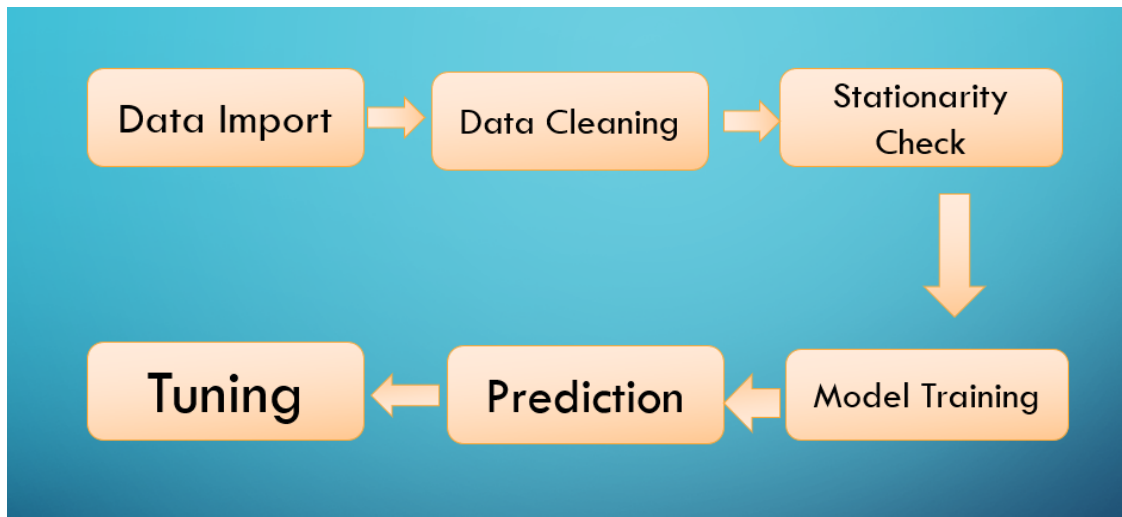
1 End to End Time Series forecasting using ARIMA

```
[25]: #Changing working directory
import os
print(os.getcwd())
import warnings
warnings.filterwarnings('ignore')
```

C:\Users\amanr\Documents\Time Series

```
[2]: from IPython.display import Image
Image(filename='Capture.png')
```

[2]:



```
[3]: #Data Import and cleaning
```

```
[26]: # Data and package Import
#Data Source - Kaggle - https://www.kaggle.com/rohanrao/nifty50-stock-market-data
import pandas as pd
import matplotlib.pyplot as plt
```

```

from statsmodels.tsa.arima_model import ARMA
TempData = pd.read_csv('HCLTECH.csv')
TempData.head(30)

```

```

[26]:
      Date  Symbol Series  Prev Close  Open  High  Low  \
0  2000-01-11  HCLTECH  EQ      580.00  1550.00  1725.00  1492.00
1  2000-01-12  HCLTECH  EQ     1554.45  1560.00  1678.85  1560.00
2  2000-01-13  HCLTECH  EQ     1678.85  1790.00  1813.20  1781.00
3  2000-01-14  HCLTECH  EQ     1813.20  1958.30  1958.30  1835.00
4  2000-01-17  HCLTECH  EQ     1958.30  2115.00  2115.00  1801.65
5  2000-01-18  HCLTECH  EQ     1801.65  1730.55  1815.00  1657.55
6  2000-01-19  HCLTECH  EQ     1774.50  1815.00  1889.00  1760.00
7  2000-01-20  HCLTECH  EQ     1851.15  1865.00  1865.00  1750.00
8  2000-01-21  HCLTECH  EQ     1757.85  1761.00  1815.00  1705.00
9  2000-01-24  HCLTECH  EQ     1781.35  1834.90  1923.90  1795.00
10 2000-01-25  HCLTECH  EQ     1923.90  1990.00  2077.85  1990.00
11 2000-01-27  HCLTECH  EQ     2077.85  2239.90  2239.90  1965.00
12 2000-01-28  HCLTECH  EQ     2080.60  2100.00  2245.00  2020.20
13 2000-01-31  HCLTECH  EQ     2196.95  2111.00  2200.00  2021.20
14 2000-02-01  HCLTECH  EQ     2030.40  1981.20  2180.00  1981.20
15 2000-02-02  HCLTECH  EQ     2162.25  2260.00  2260.00  2010.00
16 2000-02-03  HCLTECH  EQ     2024.25  2075.00  2075.00  1915.20
17 2000-02-04  HCLTECH  EQ     1956.50  1993.00  2064.00  1965.00
18 2000-02-07  HCLTECH  EQ     1986.60  1985.00  2119.95  1960.05
19 2000-02-08  HCLTECH  EQ     2062.95  2049.40  2060.00  1970.00
20 2000-02-09  HCLTECH  EQ     1999.20  2100.00  2100.00  1975.00
21 2000-02-10  HCLTECH  EQ     1999.00  1976.65  2014.80  1965.00
22 2000-02-11  HCLTECH  EQ     1996.75  1998.00  2156.50  1992.00
23 2000-02-14  HCLTECH  EQ     2156.50  2199.00  2329.05  2199.00
24 2000-02-15  HCLTECH  EQ     2324.70  2350.00  2510.70  2210.00
25 2000-02-16  HCLTECH  EQ     2508.40  2560.00  2560.00  2307.75
26 2000-02-17  HCLTECH  EQ     2423.10  2445.00  2616.95  2380.00
27 2000-02-18  HCLTECH  EQ     2614.70  2800.00  2823.90  2760.50
28 2000-02-21  HCLTECH  EQ     2775.95  2823.00  2998.00  2725.00
29 2000-02-22  HCLTECH  EQ     2819.15  2800.00  2875.00  2652.00

```

```

      Last  Close  VWAP  Volume  Turnover  Trades  \
0  1560.00  1554.45  1582.72  1192200  1.886915e+14  NaN
1  1678.85  1678.85  1657.05   344850  5.714349e+13  NaN
2  1813.20  1813.20  1804.69    53000  9.564880e+12  NaN
3  1958.30  1958.30  1939.90   270950  5.256169e+13  NaN
4  1801.65  1801.65  1990.55   428800  8.535473e+13  NaN
5  1775.00  1774.50  1716.39   359900  6.177280e+13  NaN
6  1842.80  1851.15  1842.81   316050  5.824204e+13  NaN
7  1753.50  1757.85  1801.37   204700  3.687409e+13  NaN
8  1786.00  1781.35  1774.01   282150  5.005360e+13  NaN
9  1923.90  1923.90  1875.34   328650  6.163317e+13  NaN

```

10	2077.85	2077.85	2065.71	313500	6.476010e+13	NaN
11	2078.00	2080.60	2101.92	352400	7.407150e+13	NaN
12	2190.00	2196.95	2173.00	360800	7.840197e+13	NaN
13	2021.20	2030.40	2083.45	217650	4.534620e+13	NaN
14	2160.00	2162.25	2110.66	261000	5.508834e+13	NaN
15	2035.00	2024.25	2085.91	326550	6.811542e+13	NaN
16	1965.00	1956.50	1989.66	175800	3.497821e+13	NaN
17	1977.00	1986.60	2008.25	176500	3.544560e+13	NaN
18	2036.00	2062.95	2048.51	277850	5.691791e+13	NaN
19	2014.00	1999.20	2017.09	199100	4.016027e+13	NaN
20	1991.00	1999.00	2011.50	180500	3.630763e+13	NaN
21	1990.00	1996.75	1992.36	76850	1.531132e+13	NaN
22	2156.50	2156.50	2110.61	166150	3.506781e+13	NaN
23	2329.05	2324.70	2293.27	265150	6.080618e+13	NaN
24	2500.00	2508.40	2425.21	224500	5.444586e+13	NaN
25	2418.00	2423.10	2411.48	260850	6.290343e+13	NaN
26	2616.95	2614.70	2526.71	276550	6.987609e+13	NaN
27	2771.00	2775.95	2807.40	173400	4.868032e+13	NaN
28	2805.00	2819.15	2883.95	151800	4.377830e+13	NaN
29	2710.00	2700.30	2772.79	88250	2.446984e+13	NaN

	Deliverable Volume	%Deliverble
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
5	NaN	NaN
6	NaN	NaN
7	NaN	NaN
8	NaN	NaN
9	NaN	NaN
10	NaN	NaN
11	NaN	NaN
12	NaN	NaN
13	NaN	NaN
14	NaN	NaN
15	NaN	NaN
16	NaN	NaN
17	NaN	NaN
18	NaN	NaN
19	NaN	NaN
20	NaN	NaN
21	NaN	NaN
22	NaN	NaN
23	NaN	NaN
24	NaN	NaN

25	NaN	NaN
26	NaN	NaN
27	NaN	NaN
28	NaN	NaN
29	NaN	NaN

```
[27]: #Data Cleaning
HCLTechStockData = TempData.dropna()


HCLTechStockData.index = pd.to_datetime(HCLTechStockData.Date)

HCLTechStockData = HCLTechStockData["Prev Close"]['2013-01-01':'2013-12-2']
HCLTechStockData.describe()
```

```
[27]: count      230.000000
mean       852.953478
std        156.484472
min         618.700000
25%         736.350000
50%         777.450000
75%        1023.962500
max         1161.150000
Name: Prev Close, dtype: float64
```

```
[28]: #Data Exploration
plt.figure(figsize=(16,7))
fig = plt.figure(1)
ax1 = fig.add_subplot(111)
ax1.set_xlabel('Time Frame')
ax1.set_ylabel('Stock Price for HCLTECH')
ax1.plot(HCLTechStockData)
```

```
[28]: [<matplotlib.lines.Line2D at 0x13d6eee79a0>]
```




```
[7]: # Checking stationarity
```

```
[8]: # Method 1 - Rolling Statistics
      # Method 2 - Duckkey fuller
```

```
[29]: #Determining rolling statistics
rollmean = HCLTechStockData.rolling(12).mean()
rollstd = HCLTechStockData.rolling(12).std()

plt.figure(figsize=(16,7))
fig = plt.figure(1)

#Plot rolling statistics:
orig = plt.plot(HCLTechStockData, color='blue',label='Original')
mean = plt.plot(rollmean, color='red', label='Rolling Mean')
std = plt.plot(rollstd, color='black', label = 'Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)
```



2 making Series Stationary

```
[30]: #Lets try transformation
plt.figure(figsize=(16,7))
fig = plt.figure(1)

import numpy as np
ts_log = np.log(HCLTechStockData)
plt.plot(ts_log)
```

```
[30]: [<matplotlib.lines.Line2D at 0x13d6ee15fa0>]
```




```
[31]: #Decomposition
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(ts_log,freq=1,model = 'multiplicative')

trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(16,7))
fig = plt.figure(1)

plt.subplot(411)
plt.plot(ts_log, label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
```

[31]: <matplotlib.legend.Legend at 0x13d6f496d00>




```
[32]: #Lets try differencing
plt.figure(figsize=(16,7))
fig = plt.figure(1)
ts_log_diff = ts_log - ts_log.shift()
```

```
plt.plot(ts_log_diff)

#Determining rolling statistics
rollmean = ts_log_diff.rolling(12).mean()
rollstd = ts_log_diff.rolling(12).std()


#Plot rolling statistics:
orig = plt.plot(ts_log_diff, color='blue',label='Original')
mean = plt.plot(rollmean, color='red', label='Rolling Mean')
std = plt.plot(rollstd, color='black', label = 'Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)
```



```
[13]: HCLTechStockData.sort_index(inplace= True)
```

```
[33]: from statsmodels.tsa.stattools import acf, pacf
lag_acf = acf(ts_log_diff, nlags=20)
lag_pacf = pacf(ts_log_diff, nlags=20)
```

```
[34]: import statsmodels.api as sm
fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(ts_log_diff.dropna(),lags=40,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(ts_log_diff.dropna(),lags=40,ax=ax2)
```

```
[16]: from statsmodels.tsa.arima_model import ARIMA
```

```
[47]: type(ts_log_diff)
```

```
[47]: pandas.core.series.Series
```

```
[53]: #ts_log_diff.dropna()
ts_log_diff = ts_log_diff[~ts_log_diff.isnull()]
```

```
[54]: plt.figure(figsize=(16,8))
#ts_log_diff.dropna(inplace=True)
model = ARIMA(ts_log_diff, order=(2,1,2))
results_ARIMA = model.fit()
plt.plot(ts_log_diff)
plt.plot(results_ARIMA.fittedvalues, color='red')
```

```
-----
LinAlgError                                Traceback (most recent call last)
<ipython-input-54-7d3fb1504929> in <module>
      2 #ts_log_diff.dropna(inplace=True)
      3 model = ARIMA(ts_log_diff, order=(2,1,2))
----> 4 results_ARIMA = model.fit()
      5 plt.plot(ts_log_diff)
```

```

6 plt.plot(results_ARIMA.fittedvalues, color='red')

~\anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py in fit(self,
↳ start_params, trend, method, transparams, solver, maxiter, full_output, disp,
↳ callback, start_ar_lags, **kwargs)
    1195         r, order = 'F')
    1196         """
-> 1197         mlefit = super(ARIMA, self).fit(start_params, trend,

    1198                                     method, transparams, solver,
    1199                                     maxiter, full_output, disp,

~\anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py in fit(self,
↳ start_params, trend, method, transparams, solver, maxiter, full_output, disp,
↳ callback, start_ar_lags, **kwargs)
    994         kwargs.setdefault('m', 12)
    995         kwargs.setdefault('approx_grad', True)
--> 996         mlefit = super(ARMA, self).fit(start_params, method=solver,

    997                                     maxiter=maxiter,
    998                                     full_output=full_output,
↳ disp=disp,

~\anaconda3\lib\site-packages\statsmodels\base\model.py in fit(self,
↳ start_params, method, maxiter, full_output, disp, fargs, callback, retall,
↳ skip_hessian, **kwargs)
    516         warn_convergence = kwargs.pop('warn_convergence', True)
    517         optimizer = Optimizer()
--> 518         xopt, retvals, optim_settings = optimizer._fit(f, score,
↳ start_params,

    519                                     fargs, kwargs,
    520                                     hessian=hess,

~\anaconda3\lib\site-packages\statsmodels\base\optimizer.py in _fit(self,
↳ objective, gradient, start_params, fargs, kwargs, hessian, method, maxiter,
↳ full_output, disp, callback, retall)
    213
    214         func = fit_funcs[method]
--> 215         xopt, retvals = func(objective, gradient, start_params, fargs,
↳ kwargs,

    216                                     disp=disp, maxiter=maxiter,
↳ callback=callback,

    217                                     retall=retall, full_output=full_output,

~\anaconda3\lib\site-packages\statsmodels\base\optimizer.py in _fit_lbfgs(f,
↳ score, start_params, fargs, kwargs, disp, maxiter, callback, retall,
↳ full_output, hess)
    435         func = f
    436

```

```

--> 437     retvals = optimize.fmin_l_bfgs_b(func, start_params, maxiter=maxite,
438                                     callback=callback, args=fargs,
439                                     bounds=bounds, disp=disp,

~\anaconda3\lib\site-packages\scipy\optimize\lbfgsb.py in fmin_l_bfgs_b(func,
↳ x0, fprime, args, approx_grad, bounds, m, factr, pgtol, epsilon, iprint,
↳ maxfun, maxiter, disp, callback, maxls)
195         'maxls': maxls}
196
--> 197     res = _minimize_lbfgsb(fun, x0, args=args, jac=jac, bounds=bounds,
198                             **opts)
199     d = {'grad': res['jac'],

~\anaconda3\lib\site-packages\scipy\optimize\lbfgsb.py in _minimize_lbfgsb(fun,
↳ x0, args, jac, bounds, disp, maxcor, ftol, gtol, eps, maxfun, maxiter, iprint
↳ callback, maxls, finite_diff_rel_step, **unknown_options)
358         # until the completion of the current minimization iteratio.
359         # Overwrite f and g:
--> 360         f, g = func_and_grad(x)
361         elif task_str.startswith(b'NEW_X'):
362             # new iteration

~\anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in
↳ fun_and_grad(self, x)
198         if not np.array_equal(x, self.x):
199             self._update_x_impl(x)
--> 200         self._update_fun()
201         self._update_grad()
202         return self.f, self.g

~\anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in
↳ _update_fun(self)
164     def _update_fun(self):
165         if not self.f_updated:
--> 166             self._update_fun_impl()
167             self.f_updated = True
168

~\anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in
↳ update_fun()
71
72     def update_fun():
--> 73         self.f = fun_wrapped(self.x)
74
75         self._update_fun_impl = update_fun

```

```

~\anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in
↳fun_wrapped(x)
    68         def fun_wrapped(x):
    69             self.nfev += 1
---> 70             return fun(x, *args)
    71
    72         def update_fun():

~\anaconda3\lib\site-packages\statsmodels\base\model.py in f(params, *args)
    498
    499         def f(params, *args):
--> 500             return -self.loglike(params, *args) / nobs
    501
    502         if method == 'newton':

~\anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py in loglike(self,
↳params, set_sigma2)
    808         method = self.method
    809         if method in ['mle', 'css-mle']:
--> 810             return self.loglike_kalman(params, set_sigma2)
    811         elif method == 'css':
    812             return self.loglike_css(params, set_sigma2)

~\anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py in
↳loglike_kalman(self, params, set_sigma2)
    818         Compute exact loglikelihood for ARMA(p,q) model by the Kalman
↳Filter.
    819         """
--> 820         return KalmanFilter.loglike(params, self, set_sigma2)
    821
    822         def loglike_css(self, params, set_sigma2=True):

~\anaconda3\lib\site-packages\statsmodels\tsa\kalmanf\kalmanfilter.py in
↳loglike(cls, params, arma_model, set_sigma2)
    216         paramsdtype) = cls._init_kalman_state(params, arma_model)
    217         if np.issubdtype(paramsdtype, np.float64):
--> 218             loglike, sigma2 = kalman_loglike.kalman_loglike_double(
    219
    220                 y, k, k_ar, k_ma, k_lags, int(nobs),
    221                 Z_mat, R_mat, T_mat)

statsmodels\tsa\kalmanf\kalman_loglike.pyx in statsmodels.tsa.kalmanf.
↳kalman_loglike.kalman_loglike_double()

statsmodels\tsa\kalmanf\kalman_loglike.pyx in statsmodels.tsa.kalmanf.
↳kalman_loglike.kalman_filter_double()

<__array_function__ internals> in pinv(*args, **kwargs)

```

```

~\anaconda3\lib\site-packages\numpy\linalg\linalg.py in pinv(a, rcond, hermitia)
    1959         return wrap(res)
    1960     a = a.conjugate()
-> 1961     u, s, vt = svd(a, full_matrices=False, hermitian=hermitian)
    1962
    1963     # discard small singular values

<__array_function__ internals> in svd(*args, **kwargs)

~\anaconda3\lib\site-packages\numpy\linalg\linalg.py in svd(a, full_matrices,
↳ compute_uv, hermitian)
    1624
    1625     signature = 'D->DdD' if isComplexType(t) else 'd->ddd'
-> 1626     u, s, vh = gufunc(a, signature=signature, extobj=extobj)
    1627     u = u.astype(result_t, copy=False)
    1628     s = s.astype(_realType(result_t), copy=False)

~\anaconda3\lib\site-packages\numpy\linalg\linalg.py in
↳ _raise_linalgerror_svd_nonconvergence(err, flag)
    104
    105 def _raise_linalgerror_svd_nonconvergence(err, flag):
--> 106     raise LinAlgError("SVD did not converge")
    107
    108 def _raise_linalgerror_lstsq(err, flag):

LinAlgError: SVD did not converge

```

<Figure size 1152x576 with 0 Axes>

3 Taking results back to original scale

```
[40]: ARIMA_diff_predictions = pd.Series(results_ARIMA.fittedvalues, copy=True)
print(ARIMA_diff_predictions.head())
```

```

Date
2013-01-02    0.002432
2013-01-03    0.002471
2013-01-04    0.002117
2013-01-07    0.001981
2013-01-08    0.002377
dtype: float64

```

```
[41]: ARIMA_diff_predictions_cumsum = ARIMA_diff_predictions.cumsum()
print(ARIMA_diff_predictions_cumsum.head())
```

```

Date

```

```

2013-01-02    0.002432
2013-01-03    0.004903
2013-01-04    0.007021
2013-01-07    0.009002
2013-01-08    0.011379
dtype: float64

```

```

[21]: ARIMA_log_prediction = pd.Series(ts_log.iloc[0], index=ts_log.index)
      ARIMA_log_prediction = ARIMA_log_prediction.
      ↪add(ARIMA_diff_predictions_cumsum,fill_value=0)
      ARIMA_log_prediction.head()

```

```

[21]: Date
      2013-01-01    6.427621
      2013-01-02    6.430053
      2013-01-03    6.432524
      2013-01-04    6.434641
      2013-01-07    6.436622
dtype: float64

```

```


[42]: plt.figure(figsize=(12,8))
      predictions_ARIMA = np.exp(ARIMA_log_prediction)
      plt.plot(HCLTechStockData)
      plt.plot(predictions_ARIMA)
      plt.title('RMSE: %.4f'% np.sqrt(sum((predictions_ARIMA-HCLTechStockData)**2)/
      ↪len(HCLTechStockData)))

```

```

[42]: Text(0.5, 1.0, 'RMSE: 76.7109')

```



```
[43]: results_ARIMA.predict(10,20)
```

```
[43]: Date
2013-01-15    0.001398
2013-01-16    0.001704
2013-01-17   -0.002997
2013-01-18    0.001599
2013-01-21   -0.000306
2013-01-22   -0.004052
2013-01-23   -0.001140
2013-01-24   -0.001885
2013-01-25    0.002912
2013-01-28    0.005861
2013-01-29    0.006330
dtype: float64
```

```
[24]: import pmdarima as pm
def arimamodel(timeseries):
    automodel = pm.auto_arima(timeseries,
                               start_p=3,
                               start_q=3,
                               max_p=5,
                               max_q=5,
```

```
        test="adf",
        seasonal=True,
        trace=True)

return automodel
```

```
[53]: arimamodel(ts_log)
```

Performing stepwise search to minimize aic

```
ARIMA(3,1,3)(0,0,0)[0] intercept : AIC=-1158.020, Time=0.38 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-1167.424, Time=0.03 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-1165.456, Time=0.03 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-1165.465, Time=0.07 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-1165.525, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-1163.485, Time=0.09 sec
```

Best model: ARIMA(0,1,0)(0,0,0)[0] intercept

Total fit time: 0.629 seconds

```
[53]: ARIMA(order=(0, 1, 0), scoring_args={}, suppress_warnings=True)
```

```
[ ]:
```