

ARIMA and Seasonal ARIMA

Autoregressive Integrated Moving Averages

The general process for ARIMA models is the following:

- Visualize the Time Series Data
- Make the time series data stationary
- Plot the Correlation and AutoCorrelation Charts
- Construct the ARIMA Model or Seasonal ARIMA based on the data
- Use the model to make predictions

Let's go through these steps!

```
In [3]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [4]: df=pd.read_csv('perrin-freres-monthly-champagne-.csv')
```

```
In [5]: df.head()
```

```
Out [5]:
```

	Month	Perrin Freres monthly champagne sales millions ?64-?72
0	1964-01	2815.0
1	1964-02	2672.0
2	1964-03	2755.0
3	1964-04	2721.0
4	1964-05	2946.0

```
In [6]: df.tail()
```

```
Out [6]:
```

	Month	Perrin Freres monthly champagne sales millions ?64-?72
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0
105	NaN	NaN
106	Perrin Freres monthly champagne sales millions...	NaN

```
In [7]: ## Cleaning up the data
df.columns=["Month","Sales"]
df.head()
```

```
Out [7]:
```

	Month	Sales
0	1964-01	2815.0

	Month	Sales
1	1964-02	2672.0
2	1964-03	2755.0
3	1964-04	2721.0
4	1964-05	2946.0

```
In [8]: ## Drop last 2 rows
df.drop(106,axis=0,inplace=True)
```

```
In [9]: df.tail()
```

```
Out [9]:
```

	Month	Sales
101	1972-06	5312.0
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0
105	NaN	NaN

```
In [10]: df.drop(105,axis=0,inplace=True)
```

```
In [11]: df.tail()
```

```
Out [11]:
```

	Month	Sales
100	1972-05	4618.0
101	1972-06	5312.0
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0

```
In [12]: # Convert Month into Datetime
df['Month']=pd.to_datetime(df['Month'])
```

```
In [13]: df.head()
```

```
Out [13]:
```

	Month	Sales
0	1964-01-01	2815.0
1	1964-02-01	2672.0
2	1964-03-01	2755.0
3	1964-04-01	2721.0
4	1964-05-01	2946.0

```
In [14]: df.set_index('Month',inplace=True)
```

```
In [15]: df.head()
```

```
Out [15]:
```

	Sales
Month	
1964-01-01	2815.0
1964-02-01	2672.0
1964-03-01	2755.0
1964-04-01	2721.0
1964-05-01	2946.0

```
In [16]: df.describe()
```


```
Out [16]:
```

	Sales
count	105.000000
mean	4761.152381
std	2553.502601
min	1413.000000
25%	3113.000000
50%	4217.000000
75%	5221.000000
max	13916.000000

Step 2: Visualize the Data

```
In [17]: df.plot()
```

```
Out [17]: <matplotlib.axes._subplots.AxesSubplot at 0x1d2c881a2e8>
```



```
In [18]: ### Testing For Stationarity

from statsmodels.tsa.stattools import adfuller
```

```
In [19]: test_result=adfuller(df['Sales'])
```

```
In [20]: #Ho: It is non stationary
#H1: It is stationary

def adfuller_test(sales):
    result=adfuller(sales)
    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of Observations Used']
    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis(Ho), reject the null hypothesis")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary")
```

```
In [21]: adfuller_test(df['Sales'])
```

```
ADF Test Statistic : -1.8335930563276297
p-value : 0.3639157716602417
#Lags Used : 11
Number of Observations Used : 93
weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary
```

Differencing

```
In [22]: df['Sales First Difference'] = df['Sales'] - df['Sales'].shift(1)
```

```
In [212]: df['Sales'].shift(1)
```

```
Out [212]: Month
1964-01-01      NaN
1964-02-01    2815.0
1964-03-01    2672.0
1964-04-01    2755.0
1964-05-01    2721.0
1964-06-01    2946.0
1964-07-01    3036.0
1964-08-01    2282.0
1964-09-01    2212.0
1964-10-01    2922.0
1964-11-01    4301.0
1964-12-01    5764.0
1965-01-01    7312.0
1965-02-01    2541.0
1965-03-01    2475.0
1965-04-01    3031.0
1965-05-01    3266.0
1965-06-01    3776.0
1965-07-01    3230.0
1965-08-01    3028.0
1965-09-01    1759.0
1965-10-01    3595.0
1965-11-01    4474.0
1965-12-01    6838.0
1966-01-01    8357.0
1966-02-01    3113.0
1966-03-01    3006.0
1966-04-01    4047.0
1966-05-01    3523.0
1966-06-01    3937.0
...
1970-04-01    3370.0
1970-05-01    3740.0
1970-06-01    2927.0
1970-07-01    3986.0
1970-08-01    4217.0
```

```

1970-09-01    1738.0
1970-10-01    5221.0
1970-11-01    6424.0
1970-12-01    9842.0
1971-01-01    13076.0
1971-02-01    3934.0
1971-03-01    3162.0
1971-04-01    4286.0
1971-05-01    4676.0
1971-06-01    5010.0
1971-07-01    4874.0
1971-08-01    4633.0
1971-09-01    1659.0
1971-10-01    5951.0
1971-11-01    6981.0
1971-12-01    9851.0
1972-01-01    12670.0
1972-02-01    4348.0
1972-03-01    3564.0
1972-04-01    4577.0
1972-05-01    4788.0
1972-06-01    4618.0
1972-07-01    5312.0
1972-08-01    4298.0
1972-09-01    1413.0
Name: Sales, Length: 105, dtype: float64

```

```
In [188]: df['Seasonal First Difference']=df['Sales']-df['Sales'].shift(12)
```

```
In [190]: df.head(14)
```

```
Out [190]:
```

	Sales	Sales First Difference	forecast	Seasonal First Difference
Month				
1964-01-01	2815.0	NaN	NaN	NaN
1964-02-01	2672.0	-143.0	NaN	NaN
1964-03-01	2755.0	83.0	NaN	NaN
1964-04-01	2721.0	-34.0	NaN	NaN
1964-05-01	2946.0	225.0	NaN	NaN
1964-06-01	3036.0	90.0	NaN	NaN
1964-07-01	2282.0	-754.0	NaN	NaN
1964-08-01	2212.0	-70.0	NaN	NaN
1964-09-01	2922.0	710.0	NaN	NaN
1964-10-01	4301.0	1379.0	NaN	NaN
1964-11-01	5764.0	1463.0	NaN	NaN
1964-12-01	7312.0	1548.0	NaN	NaN
1965-01-01	2541.0	-4771.0	NaN	-274.0
1965-02-01	2475.0	-66.0	NaN	-197.0

```
In [192]: ## Again test dickey fuller test
adfuller_test(df['Seasonal First Difference'].dropna())
```


```

ADF Test Statistic : -7.626619157213163
p-value : 2.060579696813685e-11
#Lags Used : 0
Number of Observations Used : 92
strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```

```
In [193]: df['Seasonal First Difference'].plot()
```

Out [193]: <matplotlib.axes._subplots.AxesSubplot at 0x1d2da817e80>




Auto Regressive Model



```
In [54]: from pandas.plotting import autocorrelation_plot
autocorrelation_plot(df['Sales'])
plt.show()
```

C:\Users\krish.naik\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel_launcher.py:2:
FutureWarning: 'pandas.plotting.autocorrelation_plot' is deprecated, import
'pandas.plotting.autocorrelation_plot' instead.




Final Thoughts on Autocorrelation and Partial Autocorrelation

- Identification of an AR model is often best done with the PACF.
 - For an AR model, the theoretical PACF “shuts off” past the order of the model. The phrase “shuts off” means that in theory the partial autocorrelations are equal to 0 beyond that point. Put another way, the number of non-zero partial autocorrelations gives the order of the AR model. By the “order of the model” we mean the most extreme lag of x that is used as a predictor.
- Identification of an MA model is often best done with the ACF rather than the PACF.
 - For an MA model, the theoretical PACF does not shut off, but instead tapers toward 0 in some manner. A clearer pattern for an MA model is in the ACF. The ACF will have non-zero autocorrelations only at lags involved in the model.

p, d, q p AR model lags d differencing q MA lags

```
In [26]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [203]: fig = plt.figure(figsize=(12,8))  
ax1 = fig.add_subplot(211)  
fig = sm.graphics.tsa.plot_acf(df['Seasonal First Difference'].iloc[13:],lags=40,ax=ax1)  
ax2 = fig.add_subplot(212)  
fig = sm.graphics.tsa.plot_pacf(df['Seasonal First Difference'].iloc[13:],lags=40,ax=ax2)
```



```
In [115]: # For non-seasonal data  
#p=1, d=1, q=0 or 1  
from statsmodels.tsa.arima_model import ARIMA
```

```
In [176]: model=ARIMA(df['Sales'],order=(1,1,1))  
model_fit=model.fit()
```

```
C:\Users\krish.naik\AppData\Local\Continuum\anaconda3\lib\site-  
packages\statsmodels\tsa\base\tsa_model.py:171: ValueWarning: No frequency information was provided, so  
inferred frequency MS will be used.  
% freq, ValueWarning)  
C:\Users\krish.naik\AppData\Local\Continuum\anaconda3\lib\site-  
packages\statsmodels\tsa\base\tsa_model.py:171: ValueWarning: No frequency information was provided, so  
inferred frequency MS will be used.  
% freq, ValueWarning)
```

```
In [177]: model_fit.summary()
```

```
Out [177]:
```

ARIMA Model Results			
Dep. Variable:	D.Sales	No. Observations:	104
Model:	ARIMA(1, 1, 1)	Log Likelihood	-951.126
Method:	css-mle	S.D. of innovations	2227.262
Date:	Wed, 18 Mar 2020	AIC	1910.251

Time: 13:40:32 BIC 1920.829
Sample: 02-01-1964 HQIC 1914.536
- 09-01-1972


	coef	std err	z	P> z	[0.025	0.975]
const	22.7838	12.405	1.837	0.069	-1.530	47.098
ar.L1.D.Sales	0.4343	0.089	4.866	0.000	0.259	0.609
ma.L1.D.Sales	-1.0000	0.026	-38.503	0.000	-1.051	-0.949

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	2.3023	+0.0000j	2.3023	0.0000
MA.1	1.0000	+0.0000j	1.0000	0.0000

```
In [178]: df['forecast']=model_fit.predict(start=90,end=103,dynamic=True)
df[['Sales','forecast']].plot(figsize=(12,8))
```

Out [178]: <matplotlib.axes._subplots.AxesSubplot at 0x1d2d9f1fda0>




```
In [113]: import statsmodels.api as sm
```

```
In [204]: model=sm.tsa.statespace.SARIMAX(df['Sales'],order=(1, 1, 1),seasonal_order=(1,1,1,12))
results=model.fit()
```

C:\Users\krish.naik\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:171: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
% freq, ValueWarning)

```
In [205]: df['forecast']=results.predict(start=90,end=103,dynamic=True)
df[['Sales','forecast']].plot(figsize=(12,8))
```

Out [205]: <matplotlib.axes._subplots.AxesSubplot at 0x1d2db70e278>



```
In [206]: from pandas.tseries.offsets import DateOffset
future_dates=[df.index[-1]+ DateOffset(months=x)for x in range(0,24)]
```

```
In [207]: future_datest_df=pd.DataFrame(index=future_dates[1:],columns=df.columns)
```

```
In [208]: future_datest_df.tail()
```


```
Out [208]:
```

	Sales	Sales First Difference	forecast	Seasonal First Difference
1974-04-01	NaN	NaN	NaN	NaN
1974-05-01	NaN	NaN	NaN	NaN
1974-06-01	NaN	NaN	NaN	NaN
1974-07-01	NaN	NaN	NaN	NaN
1974-08-01	NaN	NaN	NaN	NaN

```
In [209]: future_df=pd.concat([df,future_datest_df])
```

```
In [201]: future_df['forecast'] = results.predict(start = 104, end = 120, dynamic= True)
future_df[['Sales', 'forecast']].plot(figsize=(12, 8))
```

```
Out [201]: <matplotlib.axes._subplots.AxesSubplot at 0x1d2daee5048>
```



In []: